


Scrape Pipeline

Lesson 09 · ISBA 4715



Agenda

	Part	What
1	Scraping Concepts	What it is, etiquette, the tools we'll use (<i>slides</i>)
2	Python Pipeline	Firecrawl search + scrape, saving to <code>knowledge/raw/</code> (<i>live code</i>)
3	MCP Upgrade	Claude Code as the pipeline (<i>demo</i>)
4	Your Project	Scrape at least one source into your repo

How did Google build its first index?



Before anyone wrote an API for the web...



How did Google build its first index?



Before anyone wrote an API for the web...

Web scraping — reading pages programmatically.



Scraping is what you do when there is no API

Gathering data from websites **without an API** and **without a human clicking around in a browser**

Scrape when the data has no other door



No API exists



The API is paywalled



Content changes constantly



Automate a human task

Scrape like a guest, not a thief

Rule	Why
Check <code>robots.txt</code>	Sites declare what bots can crawl
Rate limit — <code>time.sleep()</code>	Don't hammer the server
Respect Terms of Service	Legal + ethical obligation
Expect layout changes	HTML is not a stable contract
Be mindful of PII	Don't scrape what people did not consent to share



**Scraped markdown feeds your
knowledge base**

Scraped markdown feeds your knowledge base

MP03 (API)

request → CSV → SQL warehouse

Scraped markdown feeds your knowledge base



Scraped markdown feeds your knowledge base



These feed your **knowledge base** — the unstructured side of your portfolio project.

You are building a RAG system

Retrieval-**A**ugmented **G**eneration — the standard pattern for AI systems that answer questions using your own content rather than just what a model memorized during training.

Every knowledge-base chatbot, internal-docs search, and "chat with your data" product you have seen is RAG.

LLMs made stuff up. RAG fixed that.



LLMs on their own confabulate. RAG grounds answers in documents you supply. Your project is the agentic flavor — Claude Code chooses which files in `knowledge/raw/` to read for each question.

You already have the four RAG parts



Retrieval — Firecrawl searches the web and pulls back markdown




Document store —

`knowledge/raw/` holds the scraped markdown



Indexed notes — `knowledge/wiki/`
(you build this in Milestone 02)



Generation — Claude Code reads the notes and answers grounded in them

No vector database, no embeddings — Claude Code uses grep, file reads, and its own reasoning. Simple RAG beats complex RAG until you have a reason otherwise.

One API, Search + Extract



JSON for scripts. Prose for humans.



vs.




Firecrawl (API)	WebSearch (Claude Code tool)
Returns structured data (URLs, titles, markdown)	Returns prose for the agent to read
Designed for programmatic pipelines	Built into Claude Code conversations
Same response schema every call	Different synthesized text each call
Full control: filters, domains, depth	Whatever the tool decides

Cron jobs need structured data to parse. Humans need readable prose. Pick the right tool for the reader.

We could use the SDK. We are using `requests` instead.

- Firecrawl publishes `firecrawl-py`
- You already know `requests` from MP03
- HTTP endpoints break less than SDK versions
- Same body, different wrapper

One Firecrawl call = URLs + markdown



BeautifulSoup still works. You do not need it.

- Parse HTML tag-by-tag
- Hunt for the right CSS selector
- Break every time the site redesigns



BeautifulSoup still works. You do not need it.

- Parse HTML tag-by-tag
- Hunt for the right CSS selector
- Break every time the site redesigns


This year we don't. Firecrawl does all three in one call, and handles JavaScript rendering too. If you want BeautifulSoup for interview prep, read the docs — you do not need it for this project.



Buy the API. Build the pipeline.

Problem	Don't	Do
 Find + scrape URLs	Write a crawler + parser	Firecrawl
 Click through a login	Write a bot	Playwright


Firecrawl finds. Firecrawl extracts. You save.




One API does the finding and the extracting. Your script saves the output.

MCP turns your prompt into the pipeline

What if Claude Code could call Firecrawl **directly**, without you writing Python?



MCP turns your prompt into the pipeline



MCP servers (Model Context Protocol) extend Claude Code with new tools. Install one, and Claude Code can use it inside any prompt.

Python for production. MCP for exploration.

Python (Part 02)

- ~30 lines of requests.post, loops, file writes
- Version-controlled
- Runs in GitHub Actions

MCP (Part 03)

"Find 5 URLs about Chipotle's recent earnings and save each as markdown in
knowledge/raw/."

Same result. No code.

Chipotle IR: public, rich, scrape-friendly



IR = Investor Relations



Legally required to be public



Press releases, bios, earnings



Great knowledge-base fodder

One account, 20,000 credits with STUDENTEDU



Service	URL	Free Tier
Firecrawl	firecrawl.dev	500 credits
Student Program	firecrawl.dev/student-program	20,000 credits with <code>.edu</code> email

To redeem student credits: Dashboard → **Settings** → **Billing** → apply coupon code `STUDENTEDU`.

Sign up with your `.edu` email (not GitHub OAuth) so the coupon can verify your school. No card required.

Milestone 02 needs 15 sources. Start today.



15 sources, 3+ sites



Same pattern, any URL




Job posting shapes content



Feeds your wiki pages



Same pattern as MP03, different tool



Different tools. Same shape. If you learned MP03, you already know MP04.

Next up: code it together.